

CS6533: Interactive Computer Graphics

Key Frame Animator (Part II) and Material System

10/29/2013

This week's assignment includes conceptually two separate assignments, Key Frame Animator (Part II) and Material System.

Key Frame Animator (Part II) (4 points)

In this second part of the keyframe animation system you will substitute the linear interpolation used in the first part with Catmull-Rom interpolation. This will allow for smooth animation.

The previous part of this assignment included a system for creating, deleting, editing, and reading/writing keyframes. There was also a hotkey for playing the full animation from the first keyframe to the last at small increments of time using simple linear interpolation to generate frames in between keyframes.

In this assignment you will replace in your linear interpolation code with Catmull-Rom interpolation. Remember that:

- You will need at least 4 key frames for this to work. Between 2 keyframes, you will use the one following and the one previous key frame (total of 4) to create Bezier controls using the Catmull-Rom rule.
- The above means that the animation will begin interpolating from the second keyframe and ending in the second-to-last keyframe, advancing in small time increments similarly as was done in assignment 5. This is the reason that in assignment 5, we ask you to assume the key frames to be numbered $-1, 0, \dots, n$, and the valid range for the time parameter is only $[0, n - 1]$. (Note: this is same notation we used in the Asst5 spec. In the text and notes though, we numbered the keyframes $-1, 0, \dots, n + 1$, and the valid range for the time parameter was only $[0, n]$.)
- You will have one `Cvec3` spline for each RBT's translation component and one `Quat` spline for each each RBT's rotation component.

During the animation, you will create intermediate states by evaluating the Bezier curves at intermediate times, and then use these intermediate states to display an intermediate frame. For the orientations you will need to use quaternion splines.

Specifically, you need to implement the following functions:

- In `rigtform.h`, implement the function for Catmull-Rom interpolation:

```
inline RigtForm interpolateCatmullRom(  
    const RigtForm& tform0,  
    const RigtForm& tform1,  
    const RigtForm& tform2,  
    const RigtForm& tform3,  
    const double alpha);
```

- In `quat.h`, implement the function for Catmull-Rom interpolation:

```

inline Cvec<T, n> interpolateCatmullRom(
    const Cvec<T, n>& v0,
    const Cvec<T, n>& v1,
    const Cvec<T, n>& v2,
    const Cvec<T, n>& v3,
    const double alpha);

```

- In `cvec.h`, implement the function for Catmull-Rom interpolation:

```

inline Quat interpolateCatmullRom(
    const Quat& q0,
    const Quat& q1,
    const Quat& q2,
    const Quat& q3,
    const double alpha);

```

Material System (3 points)

In this assignment we will also upgrade the scene graph infrastructure to allow for multiple types of materials with different GLSL shaders in a single scene.

Task 1: Material Infrastructure (2 points)

So far we have been using only one GLSL shader per frame during our rendering, as controlled by the global variable `g_activeShader`, an index into the global array of `ShaderState`'s `g_shaderStates[]`. A more complex graphics program, such as a modern game, will use quite a few shaders to model different surface properties. In this assignment, we will incorporate a material infrastructure that allows you to use multiple shaders in a single frame. The following sections explain the new classes that will be included in your program. Note that not until the last sub section should you start altering your code.

Uniforms

Conceptually, each shader takes in a bunch of uniform variables. Some of these uniform variables, such as the projection matrix, hold the same value for all shaders, and are set once at the beginning of the frame by you. Originally, you set it directly to the currently active GLSL shader. But if we start using a different GLSL shader in the same frame, we will need to set the values of these uniforms again, because different GLSL shaders do not know about the values of each other's uniform variables. This suggests we need some kind of data structure to hold the values of these uniform variables. Then when a new shader is selected, we populate its uniform variables with values from this repository.

In the supplied code, the `Uniforms` class defined in `uniforms.h` serves this purpose. It is essentially a dictionary mapping string name to uniform values. To use it, you would write

```

// Suppose uniforms is of type Uniforms, and m is of type Matrix4
uniforms.put("uProjection", m);

// Suppose light is of type Cvec3
uniforms.put("uLight", light);

// Set uColor variable to red
uniforms.put("uColor", Cvec3(1, 0, 0));

// You can even chain the put, since put returns the object itself
uniforms.put("a", 1)
    .put("b", 10)

```

```
.put("c", Cvec2(1, 2));
```

As you can see the first argument of `Uniforms::put` is the name of uniform variable as you have declared in the shader code. The second argument is the actual value. Right now `put` is overloaded to take the following types as the second parameter:

- A single float, int, or `Matrix4`.
- A `Cvec<T, n>` with `T` being float, int or double, and `n` being 1, 2, 3, or 4. In plain English, a int, float, or double `Cvec` with size 1, 2, 3, or 4.
- A `shared_ptr<Texture>` type.

The last bullet needs further explanation: If you recall from assignment 1, shaders can access textures by declaring a `sampler2D` typed uniform variables, e.g.,

```
uniform sampler2D uTexUnit0;
```

It can then use the GLSL function `texture()` to look up values in the texture. However, the value that you supply to this uniform variable is an integer telling it which texture unit to use, as opposed to the actual texture handle. So before calling the shader, you need to manually bind the texture to the correct texture unit first by calling the appropriate sequence of `glActiveTexture` and `glBindTexture`.

In the provided material framework, the binding of texture to texture units is automatically done behind the scene. So to provide the value of a `sampler2D` variable in a `Uniforms` object, you simply put in a shared pointer to the actual texture, encapsulated by the `Texture` class (more on this later).

Now to recap, in the original codes, you set any uniform directly by pulling the handle from `ShaderState`, and then calling one of `glUniform*` variants. In the new code, you will set the value to a `Uniforms` class, which will hold on to the value until it is needed by a shader. Thus, for the shape nodes `SgShapeNode` in the scene graph, its `draw` method now takes in a `Uniforms` object as input, instead of a `ShaderState` object. Similarly, the constructors for the `Drawer` visitor and the `Picker` visitor take in a `Uniforms` as argument, as opposed to a `ShaderState`.

RenderStates

Recall that when you draw the arcball, you enclose it with a pair of `glPolygonMode` calls to make it drawn in wireframe. OpenGL has a “state-machine” model; after you have change some states, all subsequent drawing commands use the new states. Hence after you have instructed OpenGL to draw things in wireframe, and drawn the arcball, you need to turn off the wireframe with another `glPolygonMode` call. This will quickly become hard to manage since different material properties might need different OpenGL state flags to be set, and your code will quickly become peppered with these state changing calls. (For example, in assignment 8, you will need to set states that controls how transparent objects are drawn. One of the state you will need to change for that assignment, is to enable *framebuffer blending*, by calling `glEnable(GL_BLEND)` before drawing the transparent objects. Afterward, you will need to call `glDisable(GL_BLEND)` to disable blending for non-transparent objects.)

To make your job easier, we have introduced a new class called `RenderStates`. It stores a subset of OpenGL states (that set will become larger and large as more states become relevant for our assignment). For example, it contains a `polygonMode` member function which takes the same arguments as `glPolygonMode`. The difference is that, when you set a state using `RenderStates`’s member function, that state does not immediately take effect in OpenGL, but is stored within the object. Later, when you want a `RenderStates` internal states to take effect in OpenGL, you will call its `apply()` member function. Even when you do not set a state changing member function such as `polygonMode`, the `RenderStates` has a default value for that particular states. Thus, each `RenderStates` stores the values of the same set of OpenGL states. After a `RenderStates` has been `apply`’ed, the current OpenGL states will agree with the `RenderStates`’s set of states.

So why would this be useful? Consider the following code snippets

```

// All three have a default polygonMode set to GL_FILL, and blending disabled
RenderStates r1, r2, r3;

// set r2 to be used for wireframe rendering
r2.polygonMode(GL_FRONT_AND_BACK, GL_LINE);

// set r3 to be used for transparent objects
r3.enable(GL_BLEND);

// At this point, actual OpenGL states have not been changed yet.

// Now we can switch between the three sets of render states easily

r2.apply(); // after this, GL states correspond to that of r2

// draw stuff in wireframe, and translucent;

r1.apply(); // after this, GL states correspond to that of r1

// draw stuff not in wireframe, and non translucent

r3.apply(); // after this, GL states correspond to that of r3

// draw stuff not in wireframe and translucent.

```

As you can see, there is no saving and restoring of OpenGL states. All of that are handled by the `RenderStates` internal implementation.

Moreover, a smart implementation can do lazy state setting, and perform OpenGL calls only for states that have changed between successive `apply()`.

Geometry and Texture

You have already encountered the `Geometry` and `Texture` classes before. `Geometry` encapsulates the OpenGL VBO and IBOs handles, whereas `Texture` encapsulates the OpenGL texture handles. So far we have had one type of `Geometry`: static geometry with a single vertex layout as defined by `VertexPN`. Similarly we have had a single type of `Texture`: 2D image maps. A more complex graphics program will often feature more different types of geometries and textures. For example: some geometry might need to be generated dynamically and rendered on the fly; cube texture maps (which are six 2D image textures that are mapped on an imaginary cube) might be needed to model a reflection effect; and so on.

Given that we are introducing a new Material system, and both geometry and texture need to interact with it, it is now a good time to refine our `Geometry` and `Texture` definitions a bit so that they play well with the new Material system, and can be extended to support different kind of geometries and textures in the future. The new `Geometry` and `Texture` classes are both abstract classes with certain functions that once implemented, allow custom Geometry and Texture type to be used by the Material system. You should look at the detailed comments accompanying their definition in `geometry.h` and `texture.h` for details.

To demonstrate custom types of geometry and textures, we supply one concrete implementation of `Texture` and a bunch of implementations of `Geometry`, in `texture.h` and `geometry.h`.

- **ImageTexture**: A `Texture` that loads a PPM file with three channels, and optionally stores the content in SRGB color space.
- **SimpleIndexedGeometry{PN|PNX|PNTBX}**: `Geometry` implementations that encapsulate an index buffer and a vertex buffer consisting of vertices that provide different vertex attributes

PN : Position and Normal

PNX : Position, Normal, and teXture coordinates

PNTBX : Position, Normal, Tangent, Binormal, and teXture coordinates

- **SimpleUnindexedGeometry{PN|PNX|PNTBX}**: Geometry implementations that only encapsulate a vertex buffer with vertices of different format.

The different types of geometry classes above all use a more powerful and flexible geometry implementation called **BufferObjectGeometry** under the hood. **BufferObjectGeometry** defines a geometry object using one or more vertex streams of arbitrary format, zero or one index stream, and can draw its data as different primitives (e.g., triangles, quads, points, etc.)

Material: Pulling all together

Given the above, the actual **Material** class is a light-weight container of the following:

- It keeps a shared pointer to the actual GLSL shader program used.
- It keeps a **Uniforms** field, accessible through its **getUniforms()** member function.
- It keeps a **RenderStates** field, accessible through its **getRenderStates()** member function.

The default copy semantics works fine for **Material**, so one can make a copy a **Material** by writing codes like **Material material2 = material1**. This performs a deep-copy, so if you then change the uniforms stored in **material2.getUniforms()**, it won't affect the uniforms stored in **material1**.

The **Material** has a **draw(geometry, extraUniforms)** member function that takes in a geometry and some extra uniforms. This function will first set the GLSL shader, and set all the **RenderStates** it contains. Next it binds all uniforms used by the shader, first by looking at its own **Uniforms**, and then looking at the passed in **extraUniforms**. Textures are properly bound and set at this stage as well. Then it makes sure that the supplied **geometry** can provide the vertex attributes used by the shader, and enables the corresponding vertex attribute arrays. It then calls **geometry's draw** function to draw the geometry. Finally it disables the previously enabled vertex attribute arrays.

As an example, suppose you want to use the new Material system to draw the arcball as a wireframe sphere with solid color, you would do the following

```
// Assume uniforms is of type Uniforms and already contains the projection matrix
// Assume g_arcballMat corresponds to a Material with the right renderstates and uniforms set

// Calculate the arcball's MVM

// Record it in uniforms
sendModelViewNormalMatrix(uniforms, MVM, normalMatrix(MVM));

// Draw with the material, assume g_sphere points to a Geometry for a sphere
g_arcballMat->draw(*g_sphere, uniforms);
```

The constructor of **Material** takes in two arguments, the file name of the vertex shader, and the file name of the fragment shader. And to make your life easier, if the file names contain the substring “-gl3” and the global flag **gl.GL2Compatible** is set, the corresponding shader files with “-gl2” substring will be loaded. So in your code, you only need to supply one set of file names, the ones ending with “-gl3”.

Plug into the Scene Graph

Our last step is to plug the material infrastructure into the scene graph. Also mentioned before, a shape node's **draw** function now takes in a **Uniforms** instead of **ShaderStates** as input. The same goes for the **Drawer** and **Picker** visitors.

Then our concrete implementation of the shape node, **SgGeometryShapeNode** in the replacement **scenegraph.cpp** is a simple wrapper around a shared pointer to a **Material**, a **Geometry** and the affine matrix. Its constructor now takes in a shared pointer to a full **Material**, as opposed to a simple color. Different

`SgGeometryShapeNodes` in the scene graph can use completely different materials. Its draw method will use the `Material` to draw the `Geometry` contained in the same node.

There is one catch: when we're picking, we want to use one `Material` for all shapes. To accomplish this, we keep a global variable called `g_overridingMaterial` which is a `shared_ptr<Material>`. If it is `NULL`, the `draw` function of `SgGeometryShapeNode` will do its work using its own material as usual. If `g_overridingMaterial` is not `NULL`, whatever it points to will be used by `SgGeometryShapeNode::draw()` as the material.

Finally, Code Migration

To start off, **make a copy** of your current project directory just in case you need to go back to it.

Now copy all the files in the skeleton code to your project directory, replacing files with same names in your current directory. Note that you have finished the codes in `picker.cpp`, and `scenegraph.cpp` in assignment 4. So you can either choose to copy the code you have written to the new files, or just use the newly provided files. Since what you wrote implements a well specified method, your code and the solution code should accomplish the same thing, and using either version should be fine. The `shaders` directory should also replace your project's existing `shaders` directory (Actually there're only two new shaders, `normal-*.fshader` and `normal-*.vshader`).

If you are on Mac or Linux, the new `Makefile` is already included. If you are on Windows, add all the `cpp` and `h` files to your project, **except** `asst6_5-snippets.cpp`.

Finally follow the instructions in `asst6_5-snippets.cpp` to change your `asst6.cpp` to use the new Material system. As you run the resulting program, you will notice that

- The robots are drawn with diffuse shading
- The arcball is drawn in wireframe.
- The ground is drawn with a texture.

Task 2: Make Lights Movable (1 point)

Now you need to make the two lights of the scene pickable and movable. This is a very low hanging fruit given our scene graph structure, and will help your inspect the bump mapping effect once you finish it.

Instead of storing the world position of the lights as two `Cvec3`'s `g_light1` and `g_light2`, you want to store them as Transform nodes in the scene graph. The origin of these frames serve as the position of the lights, and you can attach a Shape node to the Transform nodes so that the position of the light become visible, pickable, and manipulatable (e.g., just attach a sphere)

To pass the lights' eye space coordinates to the shader, you will need to do something like the following:

```
// get world space coordinates of the light
Cvec3 light1 = getPathAccumRbt(g_world, g_light1Node).getTranslation();

// transform to eye space, and set to uLight uniform
uniforms.put("uLight", Cvec3(invEyeRbt * Cvec4(light1, 1)));
```

assuming your `g_light1Node` points to the `SgRbtNode` corresponding to the first light.

Refer to how `g_groundNode` is set and added to `g_world` in `initScene()` on how to do this. (Don't forget to call `g_world`'s `addChild!`) You probably want to use `g_lightMat` as the material for these shape nodes.

When you're done, you should have two pickable lights in the scene. These will come in handy when we get to bump mapping in later assignments.